

Hackathon Setup Guide

Leverage featherless.ai - a serverless AI inference platform with unlimited access to thousands of open-source AI models.

Focus on building amazing applications while we handle the infrastructure.

Featherless 



Get Started in 3 Steps



Sign up for Feather Premium

- Scan the QR code or click [here](#)
- **BOA26** will automatically be entered
- Enjoy one month for free - you may need to refresh the page to see the changes

Scan here



Have fun!

Feather Premium

Active

Renews on Mar 3, 2026

\$ 25 USD / Month

Hackathon premium (100% off)

Access to DeepSeek, Kimi-K2 and GLM 4.6

Access any model - no limit on size!

Up to 4 concurrent connections

Up to 32K context

© 2026 Featherless.ai

Discord

F.A.Q.

About

Terms

Privacy

Cookies

Get Started in 3 Steps



Get your API key

→ After subscribing, go to the top right and click on API Keys

→ Here you can create an API key, keep this API key handy, you'll need it to access the service

featherless.ai Models Resources Pricing Chat Status

Search models...

FEATHERLESS

- API Keys
- Private Models
- Private Cloud

ACCOUNT

- Profile
- Change Password
- Subscription
- Payment Methods
- Referral Program

API Keys

Key	Created	Last Used
aifu rc_2e717***** *****428a	Apr 3, 2025	Oct 29, 2025
litellm rc_14a6a***** *****f182	Jun 2, 2025	Oct 16, 2025

+ New API Key


Darin Verheijke

- Account & Billing
- API Keys
- Discord
- Refer a Friend
- Log out

© 2025 Featherless.ai

Discord F.A.Q. About Terms Privacy Cookies

Get Started in 3 Steps



Choose a model

→ Go to our [model catalog](#)

→ Pick a model for your use case and copy the model ID

→ Test with a simple chat completion call!

The screenshot shows the featherless.ai website interface. The top navigation bar includes the logo, a search bar, and links for Models, Resources, Pricing, Chat, and Status. The 'Models' link is highlighted with a red box and a red arrow. Below the navigation bar, the 'Models' section is displayed, showing a grid of model cards. Each card includes a 'Warm' status indicator, a model name, a size, and a download count. The models are organized by family, with a list of families on the left side of the page.

MODEL FAMILY	Model Name	Size	Downloads
Deepseek 3	qwen25-7b-lc	(4)	8.137.502
Gemma 3	Qwen/Qwen2.5-7B-Instruct	(179)	852
Gemma 2	qwen3-0b6	(840)	755
Gemma	Qwen/Qwen3-0.6B	(13)	26
GLM 4.6	qwen25-0b5	(1)	6.456.472
GLM 4	Gensyn/Qwen2.5-0.5B-Instruct	(20)	26
GPT OSS	qwen3-4b	(1)	4.008.995
GPT-SW3	Qwen/Qwen3-4B-Instruct-2507	(5)	440
GPT2-SW3	Qwen/Qwen3-4B-Instruct-2507	(5)	440
Ling 2	Qwen/Qwen3-4B-Instruct-2507	(1)	440
Llama 3.3	Qwen/Qwen3-4B-Instruct-2507	(147)	440
Llama 3.2	Qwen/Qwen3-4B-Instruct-2507	(2505)	440
Llama 3.1	Qwen/Qwen3-4B-Instruct-2507	(1192)	440
Llama 3	Qwen/Qwen3-4B-Instruct-2507	(1188)	440
Llama 2	Qwen/Qwen3-4B-Instruct-2507	(1054)	440
Mellum	Qwen/Qwen3-4B-Instruct-2507	(6)	440
Mistral 3.1	Qwen/Qwen3-4B-Instruct-2507	(18)	440
Mistral 3	Qwen/Qwen3-4B-Instruct-2507	(120)	440

Popular Model Options



DeepSeek-V3.2

Advanced reasoning and coding capabilities



MiniMax-M2.5

Excellent in agentic tool use



Kimi-K2.5

Multimodal from the ground up



Mistral-Nemo-Instruct

Fast and efficient processing



GLM-5

Excels in long horizon tasks

Your First API Call

Using Python Requests

```
import requests
response = requests.post(
    url="https://api.featherless.ai/v1/chat/completions",
    headers={
        "Authorization": "Bearer YOUR_API_KEY"
    },
    json={
        "model": "deepseek-ai/DeepSeek-V3-0324",
        "messages": [
            {"role": "user", "content": "Hello!"}
        ]
    }
)
print(response.json())
```

Using OpenAI SDK

```
from openai import OpenAI
client = OpenAI(
    base_url="https://api.featherless.ai/v1",
    api_key="YOUR_API_KEY"
)
response = client.chat.completions.create(
    model="deepseek-ai/DeepSeek-V3-0324",
    messages=[
        {"role": "user", "content": "Hello!"}
    ]
)
print(response.choices[0].message.content)
```



Core API Endpoints

`/v1/chat/completions`

Best for: Chatbots, virtual assistants, conversational applications

- Structured messages with roles (system, user, assistant)
- Maintains conversation context automatically
- Ideal for interactive user-assistant interactions

`/v1/completions`

Best for: Content generation, text transformation, data extraction

- Takes a single prompt string
- Direct control over prompt format
- Maximum flexibility for custom use cases

Troubleshooting Common Errors



401 - Unauthenticated

API key not recognized. Verify you've copied it correctly or generate a new one from account settings.



403 - Unauthorized

Model is gated. Visit the model's page, click "Unlock Model," and agree to license terms.



500 - Internal Server Error

Request could not be processed. Check for unsupported parameters in API documentation.



503 - Service Unavailable

Insufficient capacity or cold model. Retry the request. If it persists after three attempts, report on Discord.

Concurrency Limits

Featherless operates on a unique model of **capacity reservation** rather than token-based billing. Your subscription tier determines the maximum size and number of concurrent AI model inference calls you can make, ensuring consistent performance without unpredictable token costs.

Model Concurrency Costs

7B to 15B	1	Qwen 2.5 7B, Llama2 13B
24B to 34B	2	Qwen 32B Coder, Mistral 3 24B
70B and 72B	4	Llama 3.3 70B, Qwen 2.5 72B
Deepseek v3.2, R1 & Kimi-K2.5	4	(Feather Premium only)

For a more advanced explanation and detailed examples, visit our [Concurrency Limits documentation](#).

Application Guides

Explore our detailed application guides designed to help you seamlessly integrate Featherless.ai into your existing projects and workflows.

Find these guides and more on our documentation page:

featherless.ai/docs/application-guides

Coding

Aider

Powering your in-terminal coding assistance with inference from featherless

Cursor

Setting Up Cursor with Featherless.ai

Roo Code

Roo Code (prev. Roo Cline) gives you a whole dev team of AI agents in your code editor.

Cline

A fully collaborative AI partner that's open source, fully extensible, and designed to amplify developer impact.

Workflow Automation

n8n

Supercharge your AI workflows by connecting Featherless with n8n's powerful automation platform. Build sophisticated AI-powered automations that connect with 1000+ apps and services.

Dify

Unleash 10,000+ open-source models in Dify with our dedicated plugin. Build powerful AI applications using serverless inference and low-code orchestration without the infrastructure headaches.

Ready to Build!

You now have everything you need to start building amazing AI-powered applications. Here's your action plan:

1 Test Your Setup

Get your API key and test with a simple example

2 Explore Models

Find the right model for your use case in the catalog

3 Review Examples

Check the GitHub cookbook for code samples

4 Start Building

Begin your hackathon project with confidence

Quick Links: [Sign Up](#) • [Model Catalog](#) • [Documentation](#) • [GitHub](#)

Good luck with your project! We can't wait to see what you build!